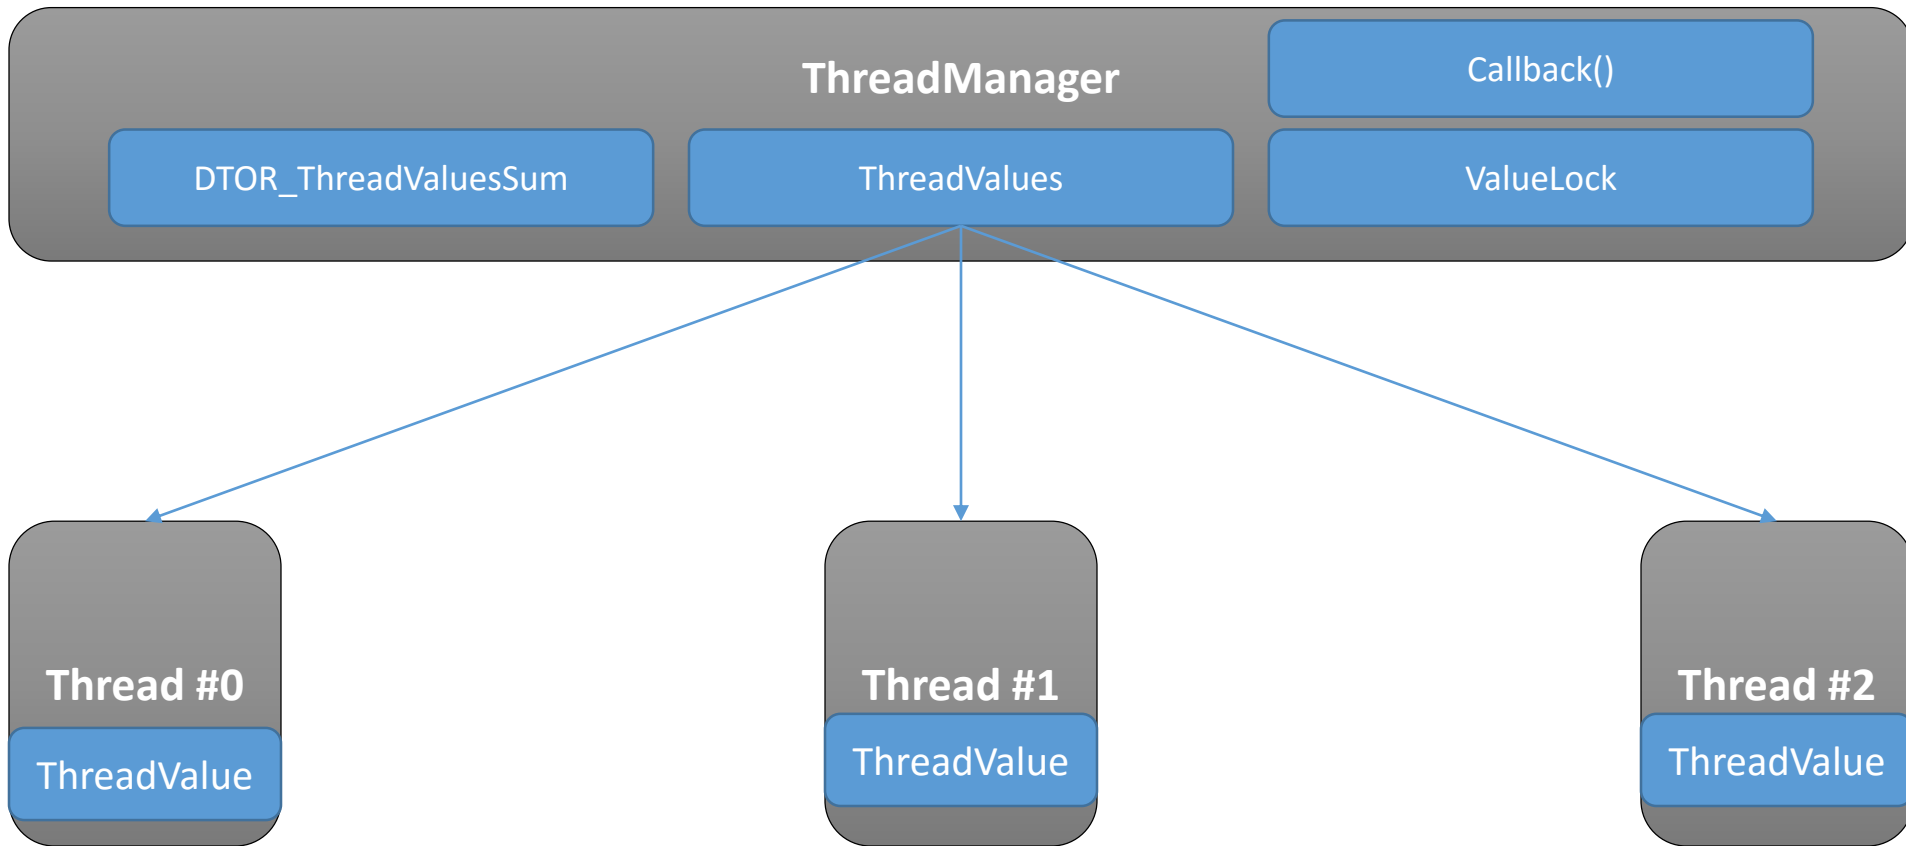


C++11 ThreadManager and ThreadValue

Managing thread-local storage from a single class

ThreadManager architecture



ThreadValue CTOR and DTOR

- `ThreadValue` is allocated as thread-local storage using C++11 `thread_local` qualifier.
- The constructor of `ThreadValue` appends the address of the latter object to the `ThreadManager::ThreadValues` vector and updates the manager, in a thread-safe way.
- The destructor of `ThreadValue` deletes the address of the latter object from the `ThreadManager::ThreadValues` vector after adding its `Value` to `ThreadManager::DTOR_ThreadValuesSum`, in a thread-safe way.

ThreadValue fields and methods

- The `ThreadValue` class contains 3 fields: `Value`, `Threshold` and a reference to a `ThreadManager`.
- `GetThreadValue()` atomically retrieves the `Value` field.
- `AdjustThreadThreshold(_X)` atomically exchanges the `Threshold` field with `Value + _X`.
- `Decrement()` decrements `Value` and then waits for the `ThreadManager::ValueLock` to be released.
- `Increment()` increments `Value`. If `Value` exceeds `Threshold`, it calls `ThreadManager::UpdateManager()` and then waits for `ThreadManager::ValueLock` to be released.

X86/AMD64 Atomicity and Memory Ordering

- The atomic Read/Write aforementioned (`ThreadValue::GetThreadValue()` and `ThreadValue::AdjustThreadThreshold()`) can of course be done using `std::atomic<>`. ➔ Reliable but not fast...
- Using Microsoft™ Visual Studio® **volatile** specific behavior, we get the desired memory ordering: (refer to <https://msdn.microsoft.com/en-us/library/12a04hfd.aspx>)

- A write to a volatile object (also known as volatile write) has Release semantics; that is, a reference to a global or static object that occurs before a write to a volatile object in the instruction sequence will occur before that volatile write in the compiled binary.
- A read of a volatile object (also known as volatile read) has Acquire semantics; that is, a reference to a global or static object that occurs after a read of volatile memory in the instruction sequence will occur after that volatile read in the compiled binary.

X86/AMD64 Atomicity and Memory Ordering

- However, the Intel® 64 and IA-32 Architectures Software Developer's Manual, Vol. 3A, « 8.1.1 Guaranteed Atomic Operations » states that:

8.1.1 Guaranteed Atomic Operations

The Intel486 processor (and newer processors since) guarantees that the following basic memory operations will always be carried out atomically:

- Reading or writing a byte
- Reading or writing a word aligned on a 16-bit boundary
- Reading or writing a doubleword aligned on a 32-bit boundary

The Pentium processor (and newer processors since) guarantees that the following additional memory operations will always be carried out atomically:

- Reading or writing a quadword aligned on a 64-bit boundary

➔ Both atomicity and ordering constraints can be obtained by using `volatile` and by aligning our data, using `__declspec(align(#))`

ThreadManager::GetGlobalValue() algorithm

In a thread-safe way:

- Acquire `ThreadManager::ValueLock`, so no further calls to `ThreadValue::Increment()/Decrement()` can occur.
- Traverse the `ThreadManager::ThreadValues` vector and compute the sum of `ThreadValue::Value`.
- Release `ThreadManager::ValueLock`.
- Return the aforementioned sum + `ThreadManager::DTOR_ThreadValuesSum`.

This retrieves the sum of thread-local values (including deleted ones) at the time the function is called, due to the lock preventing subsequent changes.

ThreadManager::UpdateManager() motivation

- The ThreadManager may need to be periodically updated if the ThreadValues (or their sum) have to be tracked.
- In this perspective, the ThreadManager class defines a user-defined Callback() function, used to set the « goal » global value for the next update, based on current global value.
 - ➔ Threshold policy: Return a new, greater threshold if the global value exceeded the old one. Otherwise, return the latter.
 - ➔ Constant policy: Return the global value + or * a given factor.
- The ThreadValue::Threshold is computed so:
 - ➔ The number of calls to ThreadManager::UpdateManager() is optimal for balanced workloads.
 - ➔ The observed global value can be greater than the « goal » global value by a ratio of MAX_ERROR at most.

ThreadManager::UpdateManager() algorithm

- In a `std::call_once` way (ie. only one thread can call the function at a given time, other threads wait and return):
 - Traverse the `ThreadManager::ThreadValues` vector and compute the sum of `ThreadValue::GetThreadValue()`.
 - `GlobalValue = ThreadManager::DTOR_ThreadValuesSum + the latter sum.`
 - `Threshold = Callback(GlobalValue)`
 - `NewMargin = std::max(Threshold - GlobalValue, Threshold * MAX_ERROR) / #Threads`


This is done to impose a minimal change, as the first expression converges to 0.

- Traverse the vector a second time, calling `ThreadValue::AdjustThreadThreshold(NewMargin)` on each element.

ThreadManager performance

- 4 threads either having their own `ThreadValue`, or sharing a single variable (either `std::atomic<>` or a `long (long)`)
- 300,000,000 Increments per thread
- Benchmarks repeated 50 times
- Configuration: Microsoft® Windows™ 10 64 bits, Visual Studio™ 2015, Release Mode x86/x64, CPU Intel® Core™ i3 3110M (2.4 GHz, Ivy Bridge, HT enabled)

Time in seconds (s)	ThreadValue (std::atomic<>)	ThreadValue (X86/AMD64)	std::atomic<>	X86/AMD64
X86 (long operands, 32b)	5.97	1.55	34.8 (5.8x)	8.69 (5.6x)
AMD64 (long long operands, 64b)	5.91	0.98	29.6 (5.0x)	8.81 (9.0x)
Due to very high contention				



ThreadManager conclusion

- Good performance due to thread locality, which can be greatly improved on X86/AMD64 platforms.
- However, it is slower and less straightforward to **retrieve the sum** of the values stored in every thread, compared to a single `std::atomic<>` shared by all threads.
 - ➔ **ThreadManager** is useful only if we do not retrieve this value often.

This program is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program. If not, see [<http://www.gnu.org/licenses/>](http://www.gnu.org/licenses/).

Intel, the Intel Logo, and Intel Core are trademarks of Intel Corporation in the U.S. and other countries.

Microsoft, Visual Studio, Windows are trademarks of Microsoft Corporation in the U.S. and other countries.